



DSC 10, Spring 2018

Lecture 26

Classification II

sites.google.com/eng.ucsd.edu/dsc-10-spring-2018

Announcements

- Lab 10 due Wednesday
 - Project 10 due Saturday
 - Please re-click download link to get updated tests
 - Please fill out CAPE
 - This is a new major and your feedback matters!
-

Classification

- Response variable is categorical; values are **classes**
 - **Binary response**: Only two classes, **0** and **1**
 - Try to **classify** the response into one of the categories, based on:
 - Values of predictor variables, called **attributes**
 - **Training set** of data in which the classes of the individuals are known
-

The Classifier (The Big Picture)

To classify a point:

- Find its k nearest neighbors
 - Take a majority vote of the k nearest neighbors to see which of the two classes appears more often
 - Assign the class that wins the majority vote
-

Finding the k Nearest Neighbors

To find the k nearest neighbors of a new point:

- Find the distance between the new point and each point in the training set
 - Augment the training data table with a column containing all the distances
 - Sort the augmented table in increasing order of the distances
 - Take the top k rows of the sorted table
-

Taking a Majority Vote

To find the class to assign the new point:

- Find the majority of the top k rows' classes
- Assign this class to the new point

Discussion Question

```
def majority(topkclasses):  
    ones = topkclasses.where('Class', are.equal_to(1)).num_rows  
    zeros = topkclasses.where('Class', are.equal_to(0)).num_rows  
    if ones > zeros:  
        return 1  
    else:  
        return 0
```

How could you implement the majority function in one line of code?

- A. `return topkclasses.group('Class', max).sort('Class', descending=True).row(0).item(1)`
- B. `return topkclasses.group('Class', max).sort('Class', descending=False).row(0).item(0)`
- C. `return topkclasses.group('Class').sort('count', descending=True).row(0).item(0)`
- D. `return topkclasses.group('Class').sort('count', descending=False).row(0).item(0)`

Measuring Accuracy

Accuracy of Classifier

- What fraction of individuals does it classify correctly?
 - Need to compare:
 - Classifier's predictions
 - True classes of individuals
 - For this, need to know the true classes. But we only know those for the training set. So now what?
-

The Test Set

- Split original training set at random into two sets
- Use one of the sets for training:
 - Explore as much as you want
 - Develop classifier
- Use the other set (test set) to compare the classifier's predictions and the true classes

Discussion Question

```
def evaluate_accuracy(training, test, k):  
    test_attributes = test.drop('Class')  
    def classify_testrow(row):  
        return classify(training, row, k)  
    c = test_attributes.apply(classify_testrow)  
    return count_equal(c, test.column('Class')) / test.num_rows
```

What is the **type** of the `test_attribute` variable?

- A. Number
- B. Array
- C. Table
- D. Row
- E. List

Discussion Question

```
def evaluate_accuracy(training, test, k):  
    test_attributes = test.drop('Class')  
    def classify_testrow(row):  
        return classify(training, row, k)  
    c = test_attributes.apply(classify_testrow)  
    return count_equal(c, test.column('Class')) / test.num_rows
```

What is the **purpose** and **return type** of the `classify_testrow` function?

- A. Predicts a class for one row, returns a number
- B. Predicts a class for the table, returns an array
- C. Predicts a class for one row, returns an array
- D. Predicts a class for the table, returns a number
- E. None of the above

Discussion Question

```
def evaluate_accuracy(training, test, k):  
    test_attributes = test.drop('Class')  
    def classify_testrow(row):  
        return classify(training, row, k)  
    c = test_attributes.apply(classify_testrow)  
    return count_equal(c, test.column('Class')) / test.num_rows
```

What is the **type** of the variable named `c`?

- A. Number
- B. Array
- C. Table
- D. Row
- E. None of the above

Discussion Question

```
def evaluate_accuracy(training, test, k):  
    test_attributes = test.drop('Class')  
    def classify_testrow(row):  
        return classify(training, row, k)  
    c = test_attributes.apply(classify_testrow)  
    return count_equal(c, test.column('Class')) / test.num_rows
```

How does this function measure [accuracy](#)?


- A. The number of 1's in the column('Class')
- B. The number of 0's in the column('Class')
- C. The number of rows where actual and predicted values are the same
- D. The proportion of rows where actual and predicted values are the same
- E. None of the above

Discussion Question

Suppose you want to test your classifier using the training set. One point becomes a *test* point and everything else is *training*. Then you repeat until each point has been the unlabeled test point once.

What value of k will give us the **largest** error (number of misclassified labels)?

- A. 0
- B. 1
- C. 5
- D. 13



Discussion Question

When we run a computer program, we'd like it to run as fast as possible. k-NN algorithm has two stages: *training* and *testing*.

Which stage will take **longer** to run: training or testing?

- A. Training
- B. Testing
- C. Same time for both
- D. Depends on the problem